Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Schärli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett (eds.), *International Conference on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA*, volume 202 of *Proceedings of Machine Learning Research*, pp. 31210–31227. PMLR, 2023. URL https://proceedings.mlr.press/v202/shi23a.html.

Karthik Valmeekam, Alberto Olmo Hernandez, Sarath Sreedharan, and Subbarao Kambhampati. Large language models still can't plan (A benchmark for llms on planning and reasoning about change). *CoRR*, abs/2206.10498, 2022. doi: 10.48550/ARXIV.2206.10498. URL https://doi.org/10.48550/arXiv.2206.10498.

Karthik Valmeekam, Kaya Stechly, and Subbarao Kambhampati. Llms still can't plan; can lrms? a preliminary evaluation of openai's o1 on planbench. 2024. URL https://api.semanticscholar.org/CorpusID:272770270.

Jason Wei, Xuezhi Wang, Dale Schuurmans, et al. Chain-of-thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022.

Gail Weiss, Yoav Goldberg, and Eran Yahav. Thinking like transformers. In Marina Meila and Tong Zhang (eds.), *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, volume 139 of *Proceedings of Machine Learning Research*, pp. 11080–11090. PMLR, 2021. URL http://proceedings.mlr.press/v139/weiss21a.html.

Tian Ye, Zicheng Xu, Yuanzhi Li, and Zeyuan Allen-Zhu. Physics of language models: Part 2.1, grade-school math and the hidden reasoning process. *arXiv preprint arXiv:2407.20311*, 2024.

Hugh Zhang, Jeff Da, Dean Lee, Vaughn Robinson, Catherine Wu, Will Song, Tiffany Zhao, Pranav Raja, Dylan Slack, Qin Lyu, Sean Hendryx, Russell Kaplan, Michele Lunati, and Summer Yue. A careful examination of large language model performance on grade school arithmetic. *CoRR*, abs/2405.00332, 2024. doi: 10.48550/ARXIV.2405.00332. URL https://doi.org/10.48550/arXiv.2405.00332.

Hattie Zhou, Arwen Bradley, Etai Littwin, Noam Razin, Omid Saremi, Joshua M. Susskind, Samy Bengio, and Preetum Nakkiran. What algorithms can transformers learn? A study in length generalization. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net, 2024. URL https://openreview.net/forum?id=AssIuHnmHX.

A Appendix

In this appendix, we provide additional details to the main text, including:

- A.1: Detailed experimental setups, including the prompt template.
- A.2: Full results on GSM8K, GSM-Symbolic, and their variants.
- A.3: Additional results for the distributional performance of several models, similar to the results from Sec. 4.1 in the main text.
- A.4: Additional results for Sec. 4.3, where we studied the impact of question difficulty. We show that fine-tuning on easier tasks does not necessarily improve performance on more difficult tasks.
- A.5: A more comprehensive discussion and analysis of performance for OpenAI o1-mini and o1-preview models.

A.1 Detailed Experimental Setup

In this work, all reported evaluations results use 8-shots with chain-of-thought prompting. We use the following prompt format:

Evaluation Prompt Format

```
// preamble or system instruction
As an expert problem solver, solve step by step the following mathematical questions.

// shot-1
Q: {{question}}
A: Let's think step by step.  {{solution}}.  The final answer is {{final answer}}.
.
.
.
// shot 8
Q: {{question}}
A: Let's think step by step.  {{solution}}.  The final answer is {{final answer}}.

// target question
Q: {{question}}
A: Let's think step by step.
```

Figure 9: The prompt format used for evaluations.

Except for the last experiment in Sec. 4.4, we use the original 8 shots from GSM8K. In addition, we allow the models to generate until either their context size limit is reached, they generate one of the end-of-response tokens such as '</s>' or '<|endoftext|>', or they finish answering the current question and move on to generating the next question, indicated by another 'Q:' generation.
Finally, we note that in all experiments we use greedy decoding to generate responses from models, with one exception: currently, the available APIs for "o1-mini" and "o1-preview" models do not allow controlling the decoding strategy, and it seems that at the time of writing, these models do not perform greedy decoding, as responses to the same prompt change.

A.2 Full Results

In Tab. 1, we present the comprehensive performance results of various models, including Gemma (Mesnard et al., 2024), Gemma2 (Rivière et al., 2024), Phi (Abdin et al., 2024), Mistral (Jiang et al., 2023), Llama3 (Dubey et al., 2024), GPT-4o (OpenAI, 2023), and the o1 (OpenAI, 2024) series, on GSM8K and its different variants, GSM-Symbolic.

We report two sets of results for GSM8K: the first column indicates the accuracy on the *full* test set of GSM8K (comprising 1,319 examples), while the second column shows the accuracy on a subset of 100 questions from the GSM8K test set, which we randomly selected to generate GSM-Symbolic templates. It is noteworthy that the performance levels across both sets are very similar, with no significant differences observed.

Table 1: Full 8-shot results of all models on GSM8Kand different variants of GSM-Symbolic.

Model	GSM8K (Full)	GSM8K (100)	Symbolic-M1	Symbolic	Symbolic-P1	Symbolic-P2	Symbolic-NoOp
Gemma2b	12.1	11.0	24.5 (± 3.85)	8.2 (± 2.21)	3.6 (± 2.13)	1.5 (± 1.63)	4.7 (± 1.99)
Gemma2b-it	12.1	11.0	16.2 (± 3.28)	8.2 (± 2.21)	1.5 (± 1.49)	1.5 (± 1.63)	4.1 (± 2.48)
Gemma-7b	53.8	50.0	34.1 (± 4.41)	25.6 (± 3.25)	26.0 (± 5.30)	3.1 (± 1.92)	8.7 (± 2.71)
Gemma-7b-it	29.3	33.0	34.1 (± 4.41)	25.6 (± 3.25)	6.0 (± 3.38)	3.1 (± 1.92)	8.7 (± 2.71)
Gemma2-2b	47.5	46.0	57.2 (± 3.40)	40.1 (± 3.04)	19.5 (± 3.89)	1.3 (± 1.37)	8.8 (± 4.12)
Gemma2-2b-it	47.5	46.0	57.2 (± 3.40)	40.1 (± 3.04)	19.5 (± 3.89)	4.5 (± 1.94)	15.7 (± 3.97)
Gemma2-9b	85.3	87.0	71.2 (± 2.81)	79.1 (± 2.99)	44.0 (± 5.69)	41.8 (± 6.00)	22.3 (± 5.11)
Gemma2-9b-it	85.3	87.0	84.4 (± 2.36)	79.1 (± 2.99)	68.1 (± 4.77)	41.8 (± 6.00)	22.3 (± 5.11)
Gemma2-27b-it	89.7	92.0	90.2 (± 1.86)	88.3 (± 2.56)	80.7 (± 4.07)	63.4 (± 4.14)	30.0 (± 3.39)
Phi-2	56.0	53.0	53.0 (± 3.10)	41.4 (± 3.56)	23.3 (± 4.07)	8.9 (± 3.33)	11.2 (± 3.51)
Phi-3-mini-128k-instruct	83.7	85.0	85.9 (± 2.44)	80.7 (± 2.94)	63.4 (± 5.63)	37.5 (± 5.76)	18.0 (± 3.83)
Phi-3-small-128k-instruct	88.5	89.0	86.4 (± 1.95)	83.7 (± 2.65)	72.0 (± 3.65)	50.7 (± 4.99)	24.5 (± 4.81)
Phi-3-medium-128k-instruct	87.3	89.0	89.6 (± 1.65)	82.5 (± 2.86)	75.8 (± 3.89)	53.1 (± 4.80)	29.4 (± 4.18)
Phi-3.5-mini-instruct	84.9	88.0	87.6 (± 1.98)	82.1 (± 3.38)	64.8 (± 5.43)	44.8 (± 6.32)	22.4 (± 4.03)
Mistral-7b-v0.1	44.5	48.0	55.4 (± 3.18)	41.1 (± 3.36)	17.4 (± 4.82)	5.5 (± 2.55)	16.2 (± 4.43)
Mistral-7b-instruct-v0.1	39.7	42.0	44.9 (± 4.29)	30.5 (± 3.47)	13.1 (± 3.51)	4.0 (± 2.24)	10.1 (± 3.42)
Mistral-7b-v0.3	40.6	44.0	54.0 (± 2.95)	40.0 (± 4.43)	15.6 (± 4.02)	3.9 (± 2.31)	16.7 (± 4.26)
Mistral-7b-instruct-v0.3	56.2	56.0	62.3 (± 2.68)	50.0 (± 3.49)	24.5 (± 4.34)	10.8 (± 3.60)	15.9 (± 4.44)
Mathstral-7b-v0.1	80.1	80.0	82.9 (± 2.87)	74.0 (± 3.49)	57.4 (± 5.20)	35.5 (± 5.07)	20.4 (± 3.58)
Llama3-8b	55.8	61.0	79.5 (± 3.62)	74.6 (± 2.94)	53.8 (± 4.54)	12.3 (± 3.43)	18.6 (± 3.86)
Llama3-8b-instruct	76.0	74.0	79.5 (± 3.62)	74.6 (± 2.94)	53.8 (± 4.54)	28.3 (± 4.37)	18.6 (± 3.86)
GPT-4o-mini	94.2	95.0	92.5 (± 1.63)	91.7 (± 2.02)	81.1 (± 3.05)	72.4 (± 4.57)	54.1 (± 3.85)
GPT-4o	95.2	95.0	94.4 (± 1.62)	94.9 (± 1.87)	93.9 (± 2.59)	88.0 (± 3.43)	63.1 (± 4.53)
o1-mini	95.1	93.0	94.9 (± 1.49)	94.5 (± 1.58)	94.3 (± 2.57)	89.1 (± 3.56)	66.0 (± 4.60)
o1-preview	94.9	96.0	93.6 (± 1.68)	92.7 (± 1.82)	95.4 (± 1.72)	94.0 (± 2.38)	77.4 (± 3.84)

A.3 Additional Results on GSM-Symbolic Performance Distributions

In section 4.1, we have presented results for several models in Fig. 2. Here, we provide additional results showing the performance on GSM-Symbolic for other models also have high variance. Moreover, these models correspond to highest drop. We also report the results for all models in table 1.

A.4 Ablation: Does Fine-Tuning on Easier Tasks Help with More Difficult Tasks?

In Sec. 4.3, we observed that the performance on GSM-P2 is significantly lower than the performance on GSM-P1. We also argued that it is unlikely that additional fine-tuning or including shots from GSM-P1 would be beneficial. Here, in Fig. 11a, we show that including shots from GSM-P1 does not improve performance compared to the results where shots come solely from GSM8K.