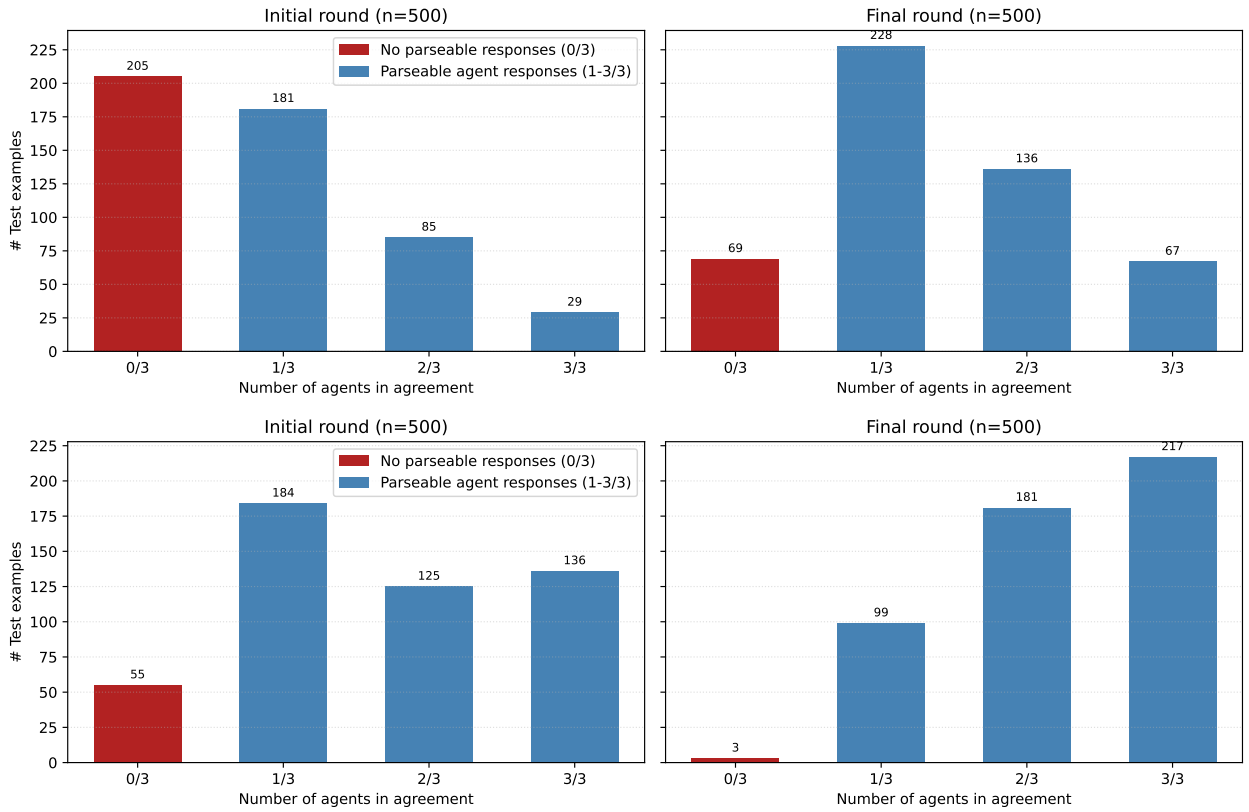

Figure 12 MACA drives meaningful improvements in both answer completeness and agent agreement, verifying the probability mass reallocation to the consensus set of reasoning trajectories (Ex: Qwen2B on GSM8K). Top: base model debate; bottom: post-trained model debate; left: initial debate round; right: final debate round.

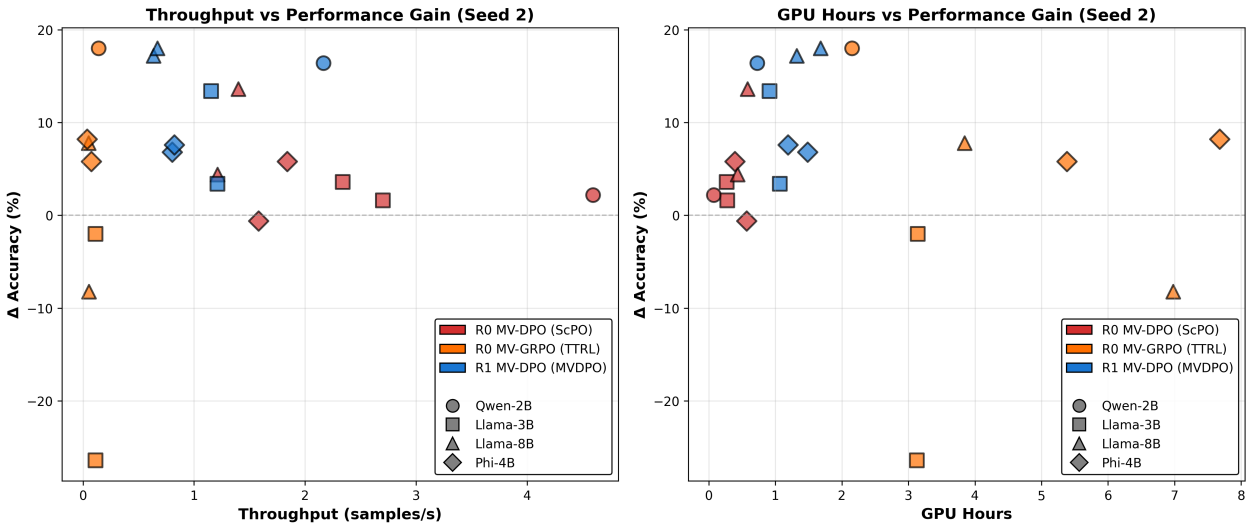

Figure 13 GPU throughput and training runtime compared with performance lift for single-agent post-training methods